

Introduction to Computer Architecture

Morteza Biglari-Abhari

Department of Electrical, Computer and Software Engineering
University of Auckland
Email: m.abhari@auckland.ac.nz

Lectures 16 - 17

Review: Our current multi-cycle implementation

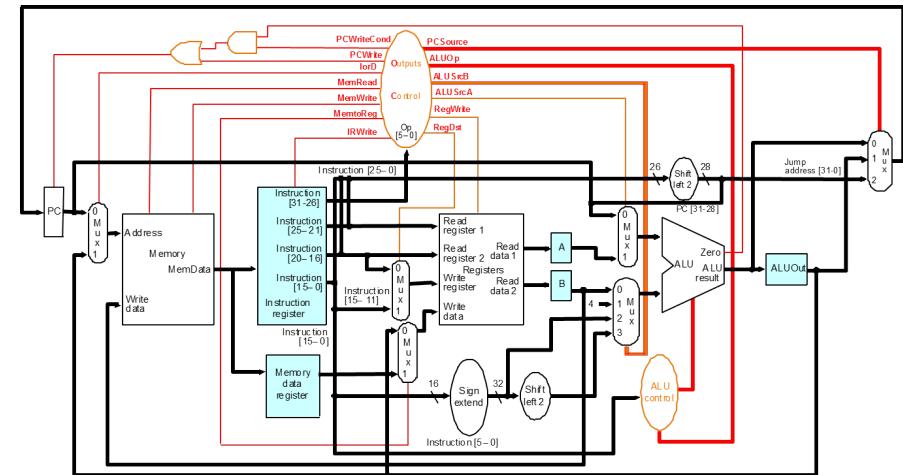


Figure 5.33 of ref3

Required clock cycles for each group of instructions

Step name	Action for R-type instructions	Action for memory reference instructions	Action for branches instructions	Action for jumps instructions
Instruction Fetch	$IR \leftarrow \text{Memory}[PC]$ $PC \leftarrow PC + 4$			
Decode/register fetch (access)	$A \leftarrow \text{Reg}[IR[25-21]]$, $B \leftarrow \text{Reg}[IR[20-16]]$ $ALUOut \leftarrow PC + (\text{sign-extend}(IR[15-0]) \ll 2)$			
Execution, address computation, branch/jump completion	$ALUOut \leftarrow A \text{ op } B$	$ALUOut = A + \text{sign-extend}(IR[15-0])$	If $(A == B)$ then $PC \leftarrow PC[31-28] \parallel (IR[25-0]) \ll 2$	
Memory access or R-type completion	$\text{Reg}[IR[15-11]] \leftarrow ALUOut$	Load: $MDR \leftarrow \text{Memory}[ALUOut]$ or Store: $\text{Memory}[ALUOut] \leftarrow B$		
Memory read completion		Load: $\text{Reg}[IR[20-16]] \leftarrow MDR$		

Adding Exception Handling to Multi-Cycle Datapath

In our simplified implementation, we only consider:

- two exceptions: **undefined instruction** and **arithmetic overflow**.
- special registers **EPC** and **Cause**
(we assume here these two registers have been added to the CPU while they are coprocessor 0 registers).

For simplicity, we assume that **Cause** register will be written **0** or **1** for **undefined instruction** and **arithmetic overflow** exceptions respectively.

When an exception occurs the content of **PC** will be changed to **0x80000180** to point to the first instruction of the **exception handler**. (This is based on the MIPS-32 architecture and might be different for different architectures).

We need three new control signals: **EPCWrite**, **CauseWrite**, **IntCause** ('1' for arithmetic overflow, and '0' for undefined instruction).
(Overflow signal from ALU is also used to signal this exception).

Our current multi-cycle implementation

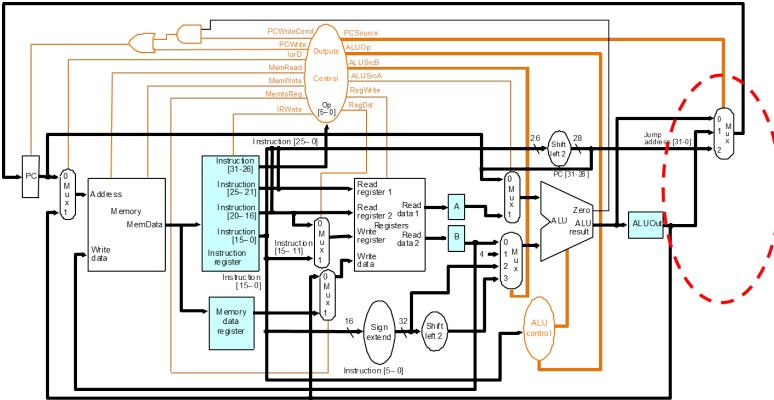


Figure 5.33 of ref3

We need to add two new internal registers **EPC** and **Cause** to our datapath in addition to some new control signals.

Signal name	Value	Effect
ALUOp	00	The ALU performs an add operation.
	01	The ALU performs a subtract operation.
	10	The funct field of the instruction determines the ALU operation.
ALUSrcB	00	The second input to the ALU comes from B register .
	01	The second input to the ALU is the constant 4 .
	10	The second input to the ALU is the sign-extended, lower 16 bits of the IR .
PCSource	00	The jump target address (IR[25-0] shifted left 2 bits) is sent to the PC .
	01	The output of the ALU (PC + 4) is sent to the PC for writing.
	10	The content of ALUOut (the branch target address) is sent to the PC for writing.
ALU control signals		
ALUop1, ALUop0		
00	Add (for loads and stores)	
01	Subtract (for beq)	
10	Depends on funct field for different R-type instructions	
11	Depends on opcode field for different ALU i-type instructions	

The multi-cycle datapath with support for exception handling

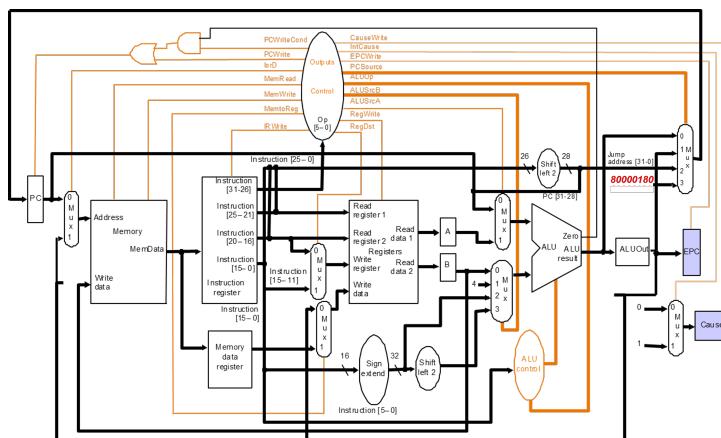
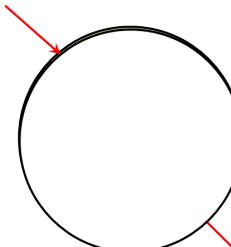


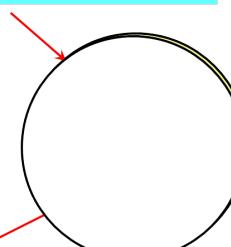
Figure 5.48 of ref3

Each state should provide control for three actions: *setting the cause register, putting the address of the instruction which caused exception into the EPC, and setting the PC to point to the exception vector address (0x80000080 in our case).*

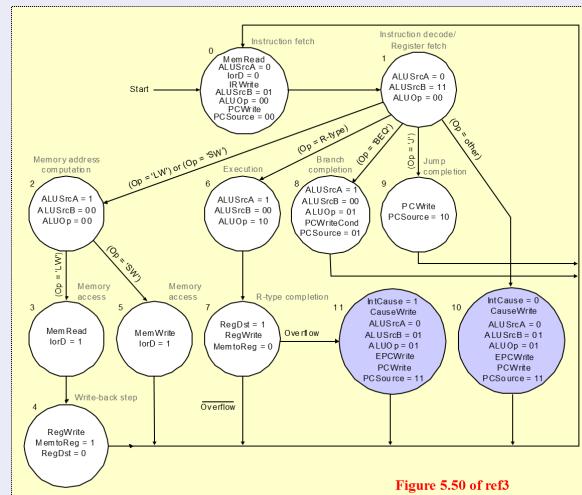
Undefined Instruction Exception



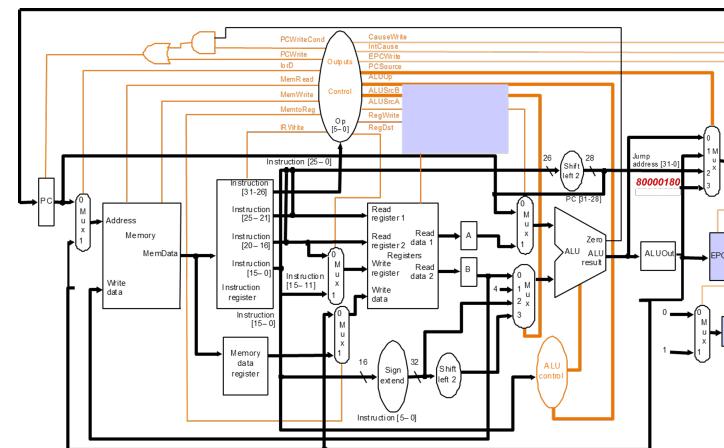
Arithmetic Overflow Exception



FSM for control unit of the multi-cycle datapath with exception detection



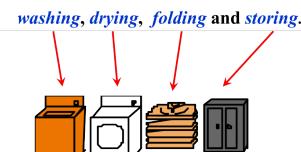
There is bug(s) in our implementation for exception handling. Can you find it?



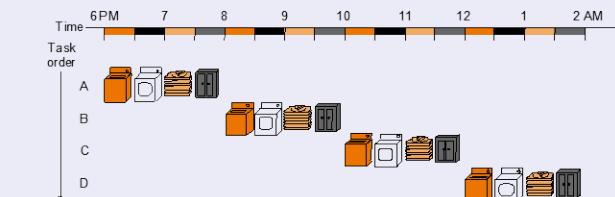
Fixing one bug in our implementation. Is there any other bug(s)?

An Overview of Pipelining (an example)

Assume 4 people (A, B, C, and D) plan to clean their clothes. The cleaning process includes:

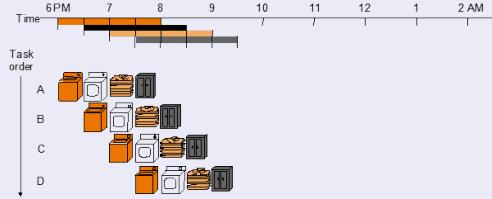


Method 1



It takes 2 hours for each person for cleaning and total time of for 4 people is 8 hours.

Method 2



It still takes 2 hours for **each person** for cleaning but the total time for 4 people is now 3.5 hours. **So pipelining improves the throughput.**

In summary

Pipelining is an implementation technique in which multiple instructions are overlapped in execution. (It does not change the execution time of a single instruction. *However, it improves the throughput*). **Pipelining is one of the key implementation techniques to make a faster CPU.**

- **Execution latency:** The time between the start and completion of an instruction.
- **Throughput:** The total amount of work done in a given time.

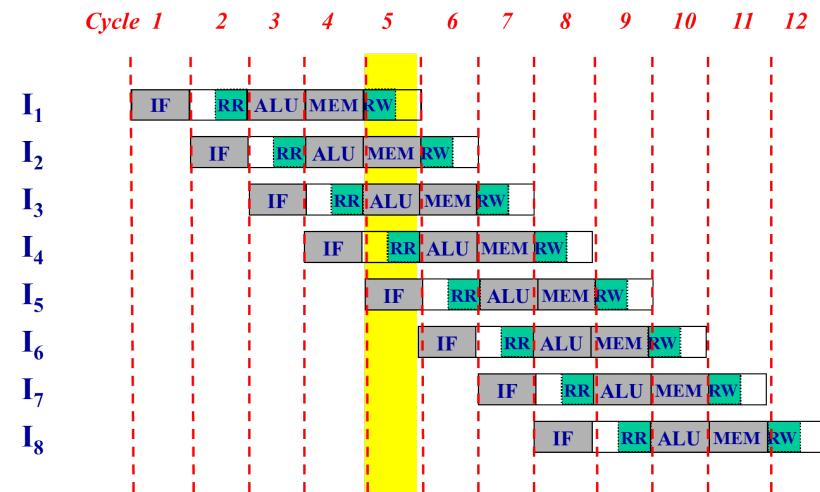
How to change the Multi-Cycle implementation to Pipelined?

R-type:	Instruction fetch	Decode & Reg. read	Execution	Reg. write (WriteBack)
Store:	Instruction fetch	Decode & Reg. read	Execution	Memory Access
Load:	Instruction fetch	Decode & Reg. read	Execution	Memory Access
Branch/ jump:	Instruction fetch	Decode & Reg. read	Execution	

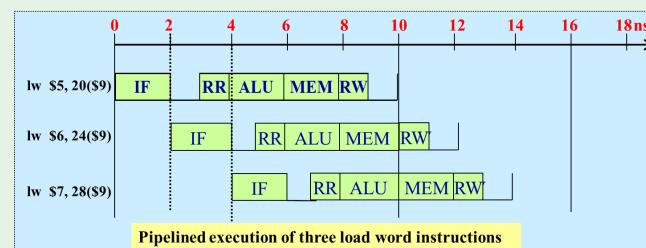
In order to change the multi-cycle implementation to pipelined:

- we notice that some resources cannot be shared as they are needed in every clock cycle for different instructions
- all instructions should have the same execution latency

Example: Pipelined execution of 8 instructions



Example



Pipelined Datapath

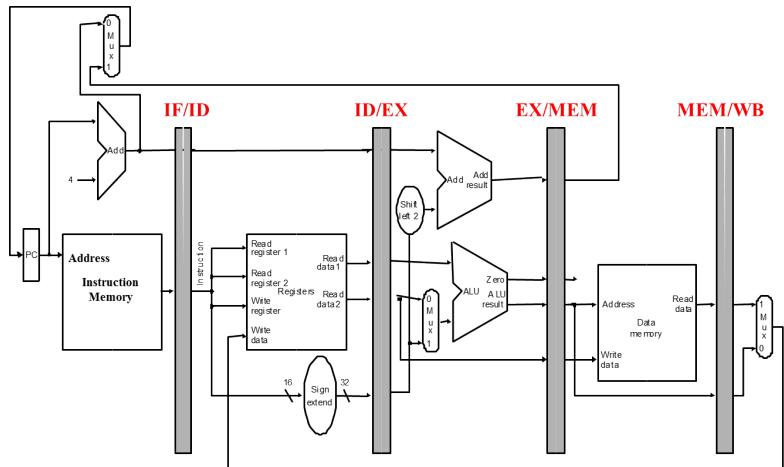


Figure 6.12 of ref3

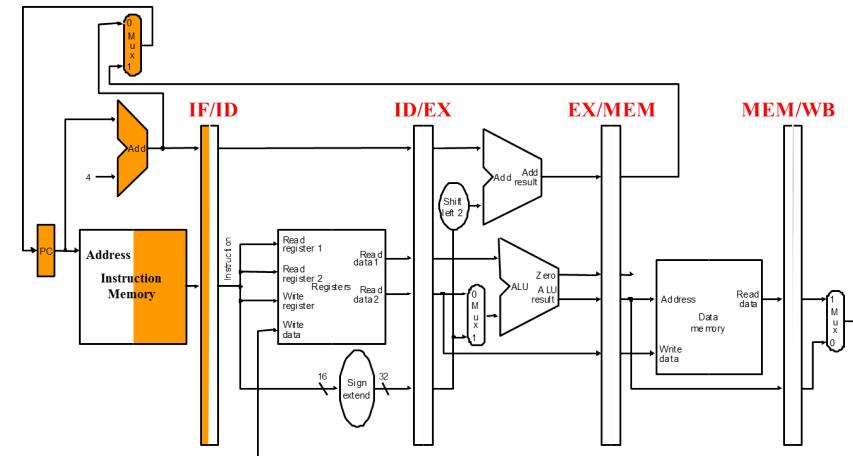


Figure 6.13 of ref3

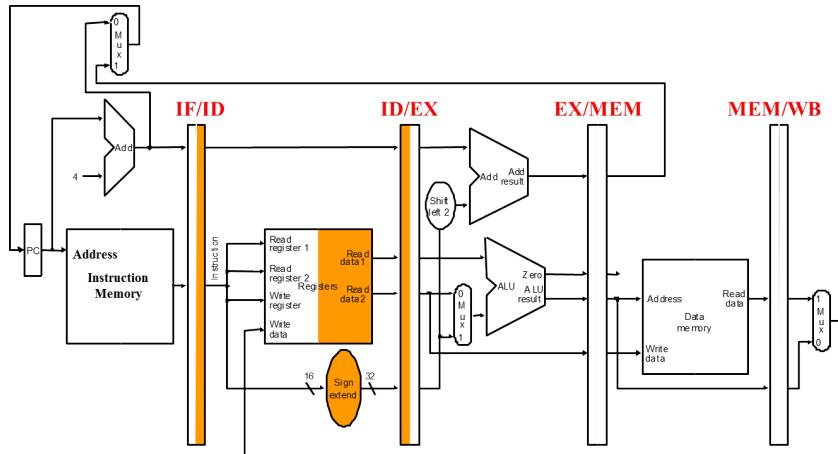


Figure 6.13 of ref3

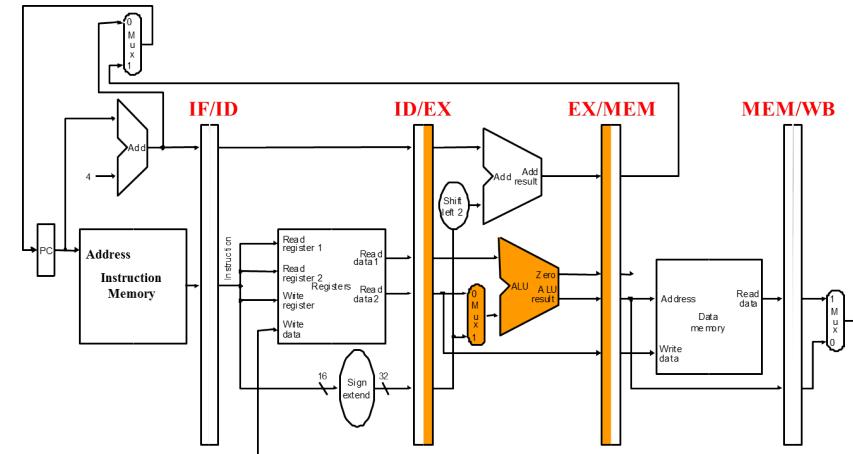


Figure 6.14 of ref3

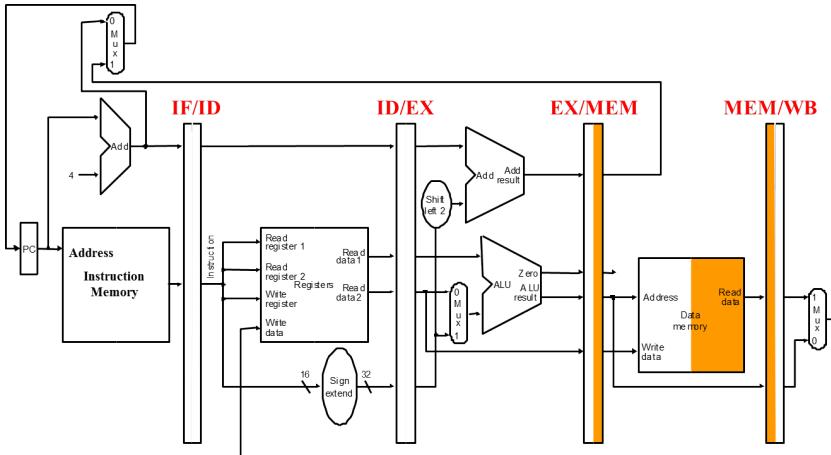


Figure 6.15 of ref3

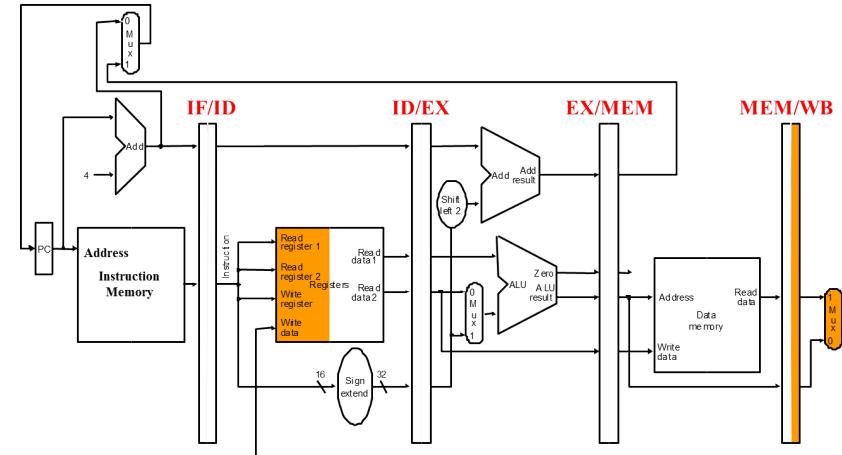


Figure 6.15 of ref3

Control Unit for the Pipelined Implementation

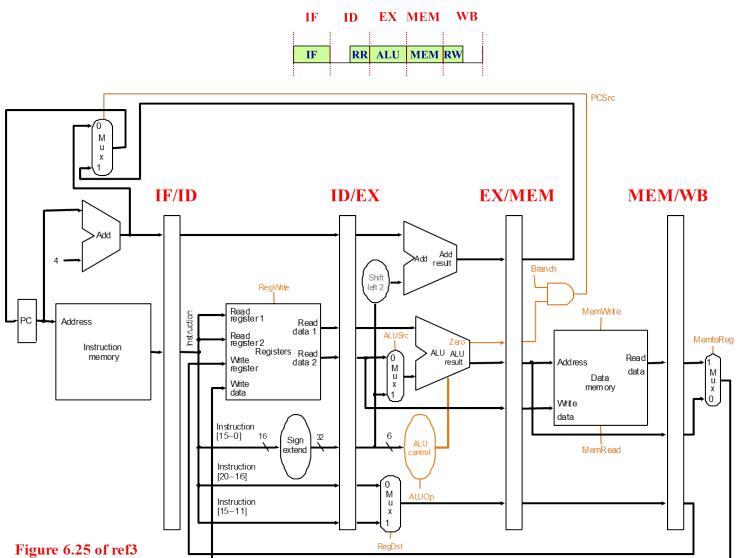


Figure 6.25 of ref3

We need to identify which control signals should be activated in each pipeline stage:

IF ID EX MEM WB



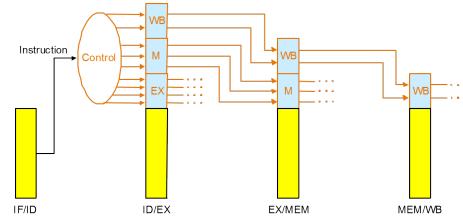
We investigate the need for a control signal in each of our following pipeline stages:

- *Instruction Fetch (IF)*
- *Instruction Decode / register file read (ID)*
- ***Execution / address calculation (EX)***
- *Memory access (MEM)*
- *Write back (WB)*

The value of control lines in each pipeline stage is given below. Note that there are no need to activate a specific control line in pipeline stages IF and ID.

Instruction	EX pipeline stage control lines				MEM pipeline stage control lines			WB pipeline stage control lines	
	Reg Dst	ALU OpI	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg Write	Memto Reg
R-type	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X

Passing the value of control lines to the appropriate points through the pipelines registers.



The pipelined datapath with its control signals

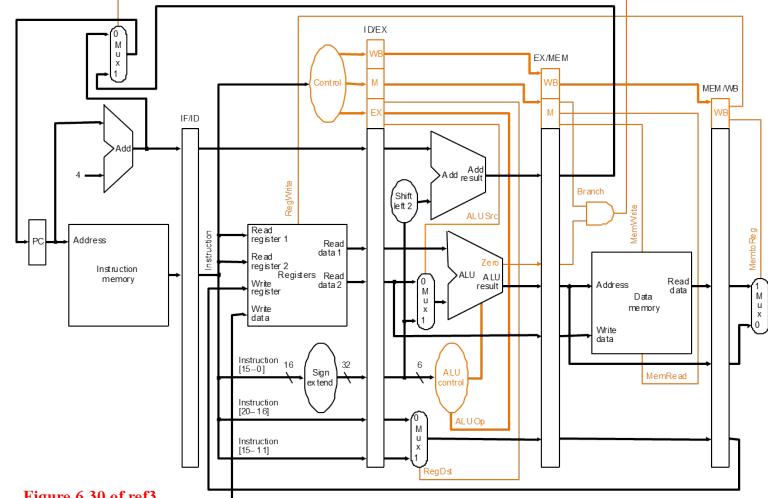
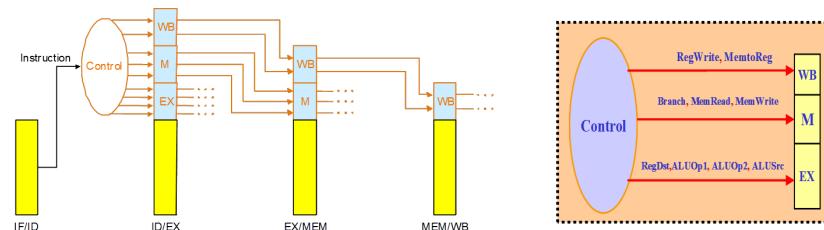


Figure 6.30 of ref3

In summary

Instruction	EX pipeline stage control lines				MEM pipeline stage control lines			WB pipeline stage control lines	
	Reg Dst	ALU OpI	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg Write	MemtoR eg
R-type	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X



Example

Using the following pipelined datapath, determine (as much as you can) the five instructions in the five pipeline stages. (If you cannot fill in a field of an instruction, state why). Use the given appendices in the following slides.

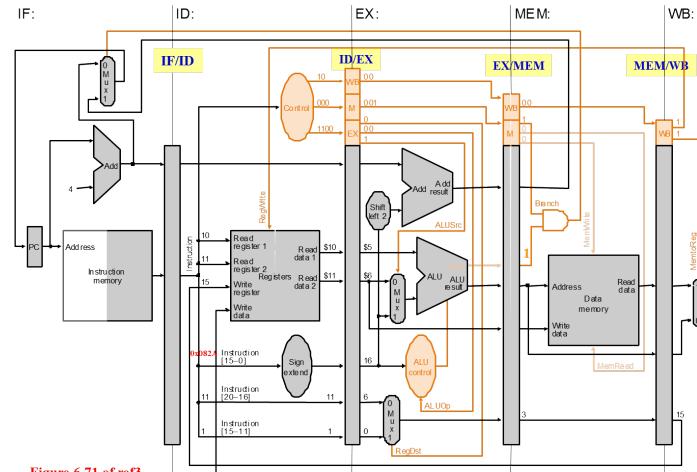


Figure 6.71 of ref3

R-type instructions	Example						
	6 bits	5 bits	5 bits	5 bits	5 bits	6 bits	
	OP	Rs	Rt	Rd	shamt	funct	
add	0	2	3	1	0	32	add \$1, \$2, \$3
addu	0	2	3	1	0	33	addu \$1, \$2, \$3
sub	0	2	3	1	0	34	sub \$1, \$2, \$3
subu	0	2	3	1	0	35	subu \$1, \$2, \$3
mfc0	16	0	1	14	0	0	mfc0 \$1, \$epc
mult	0	2	3	0	0	24	mult \$2, \$3
multu	0	2	3	0	0	25	multu \$2, \$3
div	0	2	3	0	0	26	div \$2, \$3
divu	0	2	3	0	0	27	divu \$2, \$3
mfhi	0	0	0	1	0	16	mfhi \$1
mflo	0	0	0	1	0	18	mflo \$1
and	0	2	3	1	0	36	and \$1, \$2, \$3
or	0	2	3	1	0	37	or \$1, \$2, \$3

R-type instructions (Cnt.)	Example						All values in decimal
	6 bits	5 bits	5 bits	5 bits	5 bits	6 bits	
OP	Rs	Rt	Rd	shamt	funct		
sll	0	0	2	1	10	0	sll \$1, \$2, 10
srl	0	0	2	1	10	2	srl \$1, \$2, 10
slt	0	2	3	1	0	42	slt \$1, \$2, \$3
sltu	0	2	3	1	0	43	sltu \$1, \$2, \$3
jr	16	0	1	14	0	0	jr \$31

J-type instructions	Example						
	6 bits	5 bits	5 bits	5 bits	5 bits	6 bits	
	OP	target address					
j	2	2500					j 10000
jal	3	2500					jal 10000

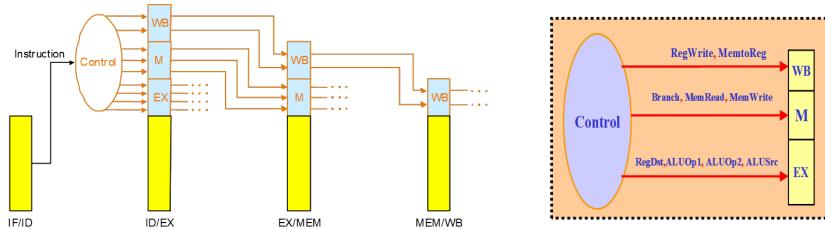
i-type instructions	Example						
	6 bits	5 bits	5 bits	5 bits	5 bits	6 bits	
	OP	Rs	Rt	Immediate/offset			
addi	8	2	1	100			addi \$1, \$2, 100
addiu	9	2	1	100			addiu \$1, \$2, 100
andi	12	2	1	100			andi \$1, \$2, 100
ori	13	2	1	100			ori \$1, \$2, 100
lw	35	2	1	100			lw \$1, 100(\$2)
sw	43	2	1	100			sw \$1, 100(\$2)
lui	15	0	1	100			lui \$1, 100
beq	4	1	2	100			beq \$1, \$2, 100
bne	5	1	2	100			bne \$1, \$2, 100
slti	10	2	1	100			slti \$1, \$2, 100
sltiu	11	2	1	100			sltiu \$1, \$2, 100

Instruction Opcode	ALUop	Instruction Operation	Funct field Inst[5-0]	Desired ALU action	ALU control input
LW	00	load word	xxxxxx	add	010
SW	00	store word	xxxxxx	add	010
BEQ	01	branch equal	xxxxxx	subtract	110
R-type	10	add	100000	add	010
R-type	10	subtract	100010	subtract	110
R-type	10	AND	100100	and	000
R-type	10	OR	100101	or	001
R-type	10	Set on less than	101010	set on less than	111

How the ALU control bits are set depends on the ALUop control bits and the different function codes for the R-type instruction.

Figure 5.14 of ref3

Instruction	EX pipeline stage control lines				MEM pipeline stage control lines			WB pipeline stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg Write	MemtoR eg
R-type	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X



IF stage:

ID stage:

EX stage:

Example

Consider a non-pipelined processor (machine M1), which has a 1 ns clock cycle period. It takes 4 cycles for ALU operations and branches, and 5 cycles for memory operations to execute. Assume that the relative frequencies of these operations are 40%, 20%, 40% respectively. Suppose we designed a pipelined processor (machine M2), which has a 1.2 ns clock cycle period (0.2 ns is added to 1 ns non-pipelined cycle due to pipeline implementation issues such as clock skew, setup/hold times, ...).

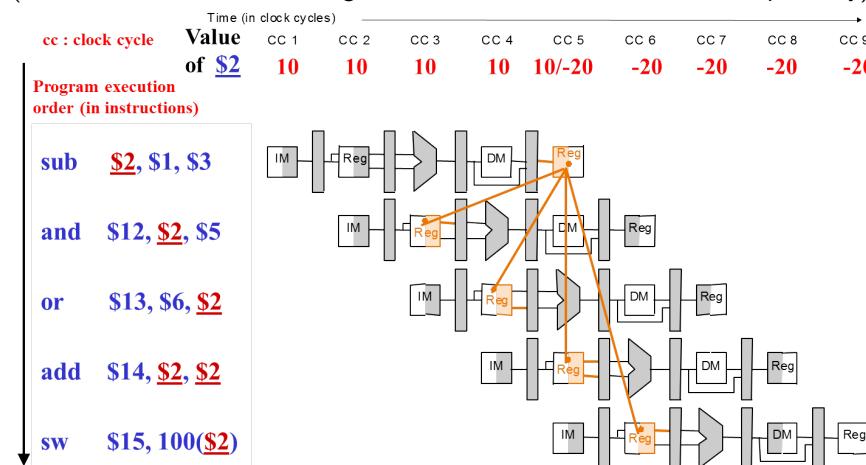
How much speedup can be achieved through using machine M2?

MEM stage:

WB stage:

Pipeline Hazards

Are the results of the following code correct (using our pipelined implementation)?
(Assume the initial values in regs. \$1, \$2 and \$3 are 10, 10 and 30 respectively).



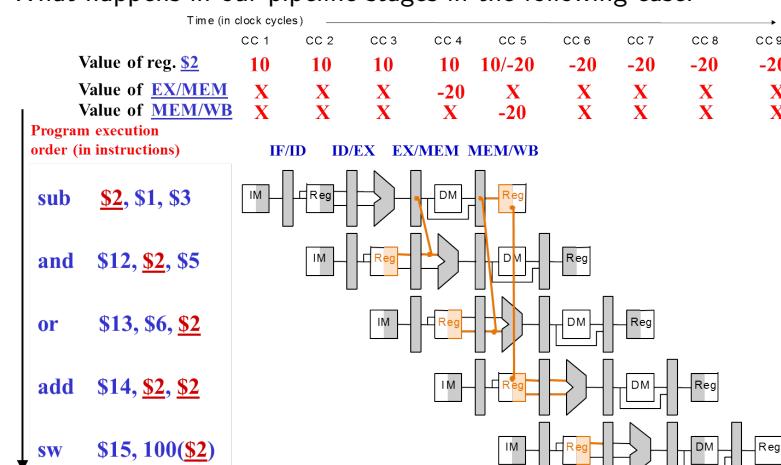
There are situations, called ***pipeline hazards***, that prevent the next instruction in the instruction stream to execute properly during its designated clock cycle. **Hazards** reduce the performance and if not handled properly wrong results may be generated. So, pipeline might be **stalled** or some other ways should be used to prevent wrong results.

Three different types of pipeline hazards:

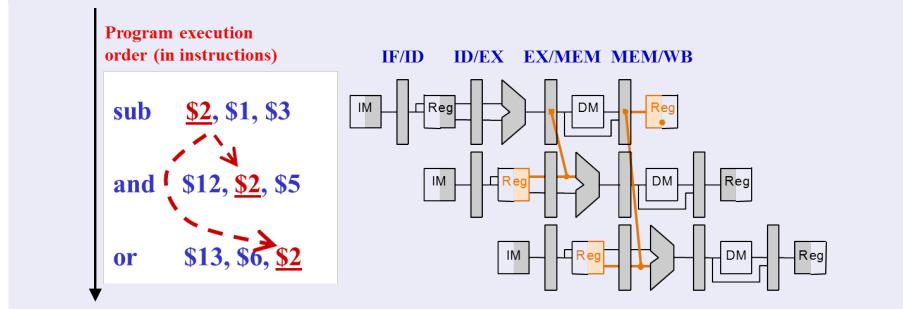
- **Structural hazards**, which are present when resource conflicts occur when instructions simultaneously access the same resource.
- **Data hazards**, which occur when an instruction depends on the results of a previous instruction which is needed to be ready for the execution of the current instruction.
- **Control hazards** may occur when branches are in the pipeline with other instructions and the PC value is to be changed.

Detecting and resolving Data Hazards in the pipeline

What happens in our pipeline stages in the following case:



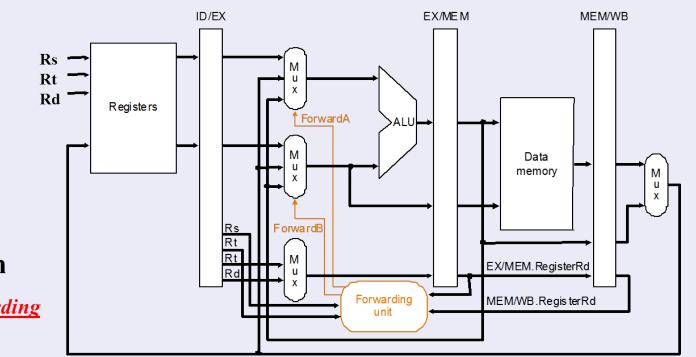
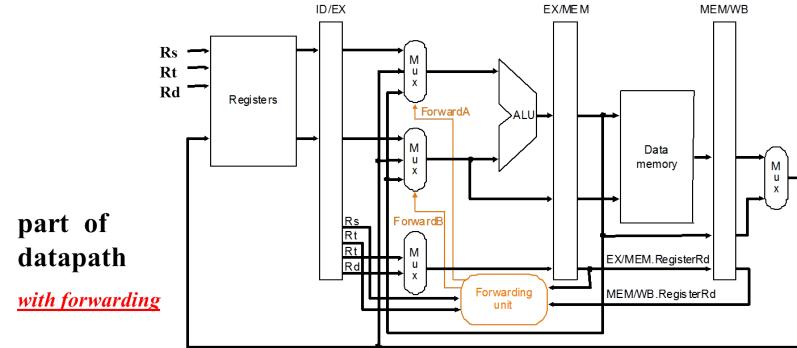
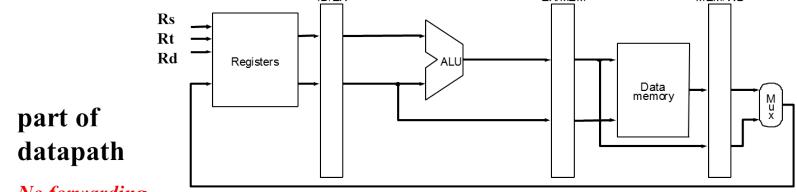
How to detect data hazards?



Data hazards can be detected through checking the proper pipeline registers.



The above is valid for ALU instructions. So, **RegWrite** control signal should be checked.

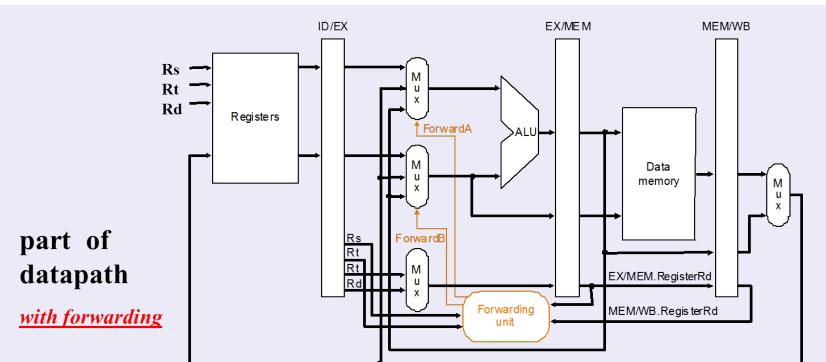


Detecting
data hazard
in EX stage
and
forwarding:

```

if (EX/MEM.RegWrite & (EX/MEM.RegisterRd != 0)
& (EX/MEM.RegisterRd = ID/EX.RegisterRs))
    ForwardA = 10

if (EX/MEM.RegWrite & (EX/MEM.RegisterRd != 0)
& (EX/MEM.RegisterRd = ID/EX.RegisterRt))
    ForwardB = 10
  
```



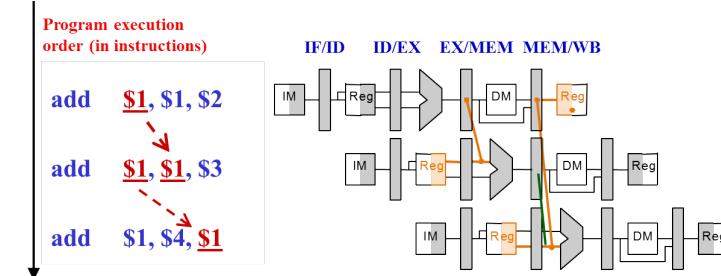
Detecting
data hazard
in MEM
stage and
forwarding:

```

if (MEM/WB.RegWrite & (MEM/WB.RegisterRd != 0)
& (MEM/WB.RegisterRd = ID/EX.RegisterRs)
& (EX/MEM.RegisterRd != ID/EX.RegisterRs))
    ForwardA = 01

if (MEM/WB.RegWrite & (MEM/WB.RegisterRd != 0)
& (MEM/WB.RegisterRd = ID/EX.RegisterRt)
& (EX/MEM.RegisterRd != ID/EX.RegisterRt))
    ForwardB = 01
  
```

Summary of forwarding mechanism (for Data Hazard):



The above data hazard can be resolved through result **forwarding** (for ALU instructions).

- We check at the EX stage if one of the source registers of the instruction in ID stage is the same as the destination register of the previous instruction(s) either in the EX/MEM or MEM/WB pipeline registers. If that is the case, then the appropriate value is forwarded from the pipeline registers (EX/MEM or MEM/WB) to the input of ALU.

Control values for the forwarding multiplexors:

Mux Control	Source	Explanation
ForwardA = 00	ID/EX	The first ALU operand comes from the register file.
ForwardA = 10	EX/MEM	The first ALU operand is forwarded from the prior ALU result.
ForwardA = 01	MEM/WB	The first ALU operand is forwarded from data memory or an earlier ALU result.
ForwardB = 00	ID/EX	The second ALU operand comes from the register file.
ForwardB = 10	EX/MEM	The second ALU operand is forwarded from the prior ALU result.
ForwardB = 01	MEM/WB	The second ALU operand is forwarded from data memory or an earlier ALU result.

References and Recommended Readings

- ① David A. Patterson and John L. Hennessy, *Computer Organization & Design – The Hardware/Software Interface*, 4th Edition Morgan Kaufmann Publishers 2009. ([Chapters 1 to 4](#)) and [Appendix A \(HP_AppA.pdf\)](#)
- ② John L. Hennessy and David A. Patterson, *Computer Architecture: A Quantitative Approach*, 3rd Edition, Morgan Kaufmann Publishers (Elsevier Science) 2003. ([Chapters 2](#))
- ③ David A. Patterson and John L. Hennessy, *Computer Organization & Design – The Hardware/Software Interface*, 2nd Edition Morgan Kaufmann Publishers 1998. ([Chapters 4, 5 and 6](#))